

Water Potability Prediction



The Battleship Team



Final Project

Our Team



**Muhammad
Nurfalah
Rohmawan**



Sandy Febrian



Samsul Rizal



Renaldi Ega

Table of Content

- A. Background of Study
- B. Materials and Methods (EDA, Data Processing, Modelling, etc)
- C. Conclusion



A. Background of Study

- Human existence relies on water, but both surface and groundwater sources are often polluted by various contaminants generated from rapid population growth and urban development.
- Globally, 20% of people lack clean drinking water, and 50% lack safe sanitation systems, as per the United Nations Environment Program (2000).
- In recent years, predicting water quality has become a significant research area due to its critical importance to address any issue related to clean water in the future.
- Predictive models can provide early warnings for water quality issues, such as contamination events or natural disasters. This information is invaluable for emergency response and safeguarding public health.



B. Material and Method

In this project, we aim to predict the water potability based on a set of features. To train our machine learning model, we will use Scikit Learn and deploy to Streamlit

Source: <https://www.kaggle.com/datasets/uom190346a/water-quality-and-potability/data>



Evaluating Data Structure

1 df

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379783	86.990970	2.963135	0
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.180013	56.329076	4.500656	0
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.868637	66.420093	3.055934	0
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436524	100.341674	4.628771	0
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558279	31.997993	4.075075	0
...
3271	4.668102	193.681735	47580.991603	7.166639	359.948574	526.424171	13.894419	66.687695	4.435821	1
3272	7.808856	193.553212	17329.802160	8.061362	NaN	392.449580	19.903225	NaN	2.798243	1
3273	9.419510	175.762646	33155.578218	7.350233	NaN	432.044783	11.039070	69.845400	3.298875	1
3274	5.126763	230.603758	11983.869376	6.303357	NaN	402.883113	11.168946	77.488213	4.708658	1
3275	7.874671	195.102299	17404.177061	7.509306	NaN	327.459760	16.140368	78.698446	2.309149	1

3276 rows x 10 columns

Attributes Information

pH: The pH level of the water.

Hardness: Water hardness, a measure of mineral content.

Solids: Total dissolved solids in the water.

Chloramines: Chloramines concentration in the water.

Sulfate: Sulfate concentration in the water.

Conductivity: Electrical conductivity of the water.

Organic_carbon: Organic carbon content in the water.

Trihalomethanes: Trihalomethanes concentration in the water.

Turbidity: Turbidity level, a measure of water clarity.

Potability: Target variable; indicates water potability with values 1 (potable) and 0 (not potable).



Evaluating Data Structure

```
[ ] 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3276 entries, 0 to 3275  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype    
---  ---                  
0    ph                    2785 non-null   float64  
1    Hardness              3276 non-null   float64  
2    Solids                3276 non-null   float64  
3    Chloramines           3276 non-null   float64  
4    Sulfate               2495 non-null   float64  
5    Conductivity          3276 non-null   float64  
6    Organic_carbon        3276 non-null   float64  
7    Trihalomethanes       3114 non-null   float64  
8    Turbidity             3276 non-null   float64  
9    Potability            3276 non-null   int64    
dtypes: float64(9), int64(1)  
memory usage: 256.1 KB
```




Evaluating Data Structure

```
✓ [6] 1 df.describe()
```

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes	Turbidity	Potability
count	2785.000000	3276.000000	3276.000000	3276.000000	2495.000000	3276.000000	3276.000000	3114.000000	3276.000000	3276.000000
mean	7.080795	196.369496	22014.092526	7.122277	333.775777	426.205111	14.284970	66.396293	3.966786	0.390110
std	1.594320	32.879761	8768.570828	1.583085	41.416840	80.824064	3.308162	16.175008	0.780382	0.487849
min	0.000000	47.432000	320.942611	0.352000	129.000000	181.483754	2.200000	0.738000	1.450000	0.000000
25%	6.093092	176.850538	15666.690297	6.127421	307.699498	365.734414	12.065801	55.844536	3.439711	0.000000
50%	7.036752	196.967627	20927.833607	7.130299	333.073546	421.884968	14.218338	66.622485	3.955028	0.000000
75%	8.062066	216.667456	27332.762127	8.114887	359.950170	481.792304	16.557652	77.337473	4.500320	1.000000
max	14.000000	323.124000	61227.196008	13.127000	481.030642	753.342620	28.300000	124.000000	6.739000	1.000000

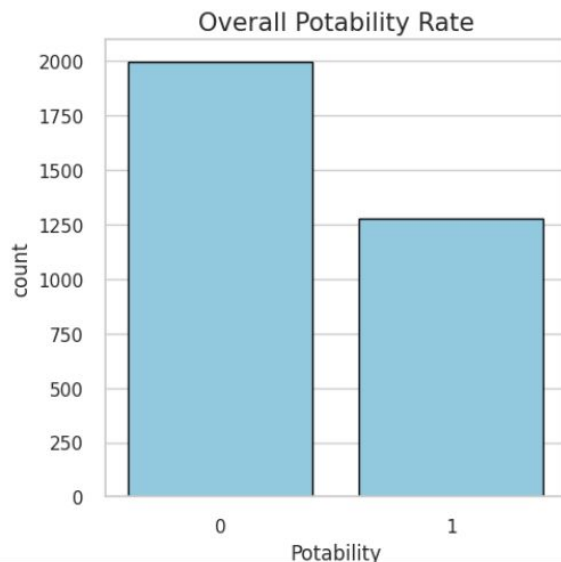
```
✓ [7] 1 df.isna().sum()
```

```
ph          491
Hardness    0
Solids      0
Chloramines 0
Sulfate     781
Conductivity 0
Organic_carbon 0
Trihalomethanes 162
Turbidity   0
Potability  0
dtype: int64
```













Checking Target Feature

```
1 # Check target feature
2
3 plt.figure(figsize=(5,5))
4 plt.title('Overall Potability Rate', fontsize=15)
5 sns.countplot(data=df, x=df['Potability'], color='skyblue', edgecolor='black')
6
7 plt.show()
```



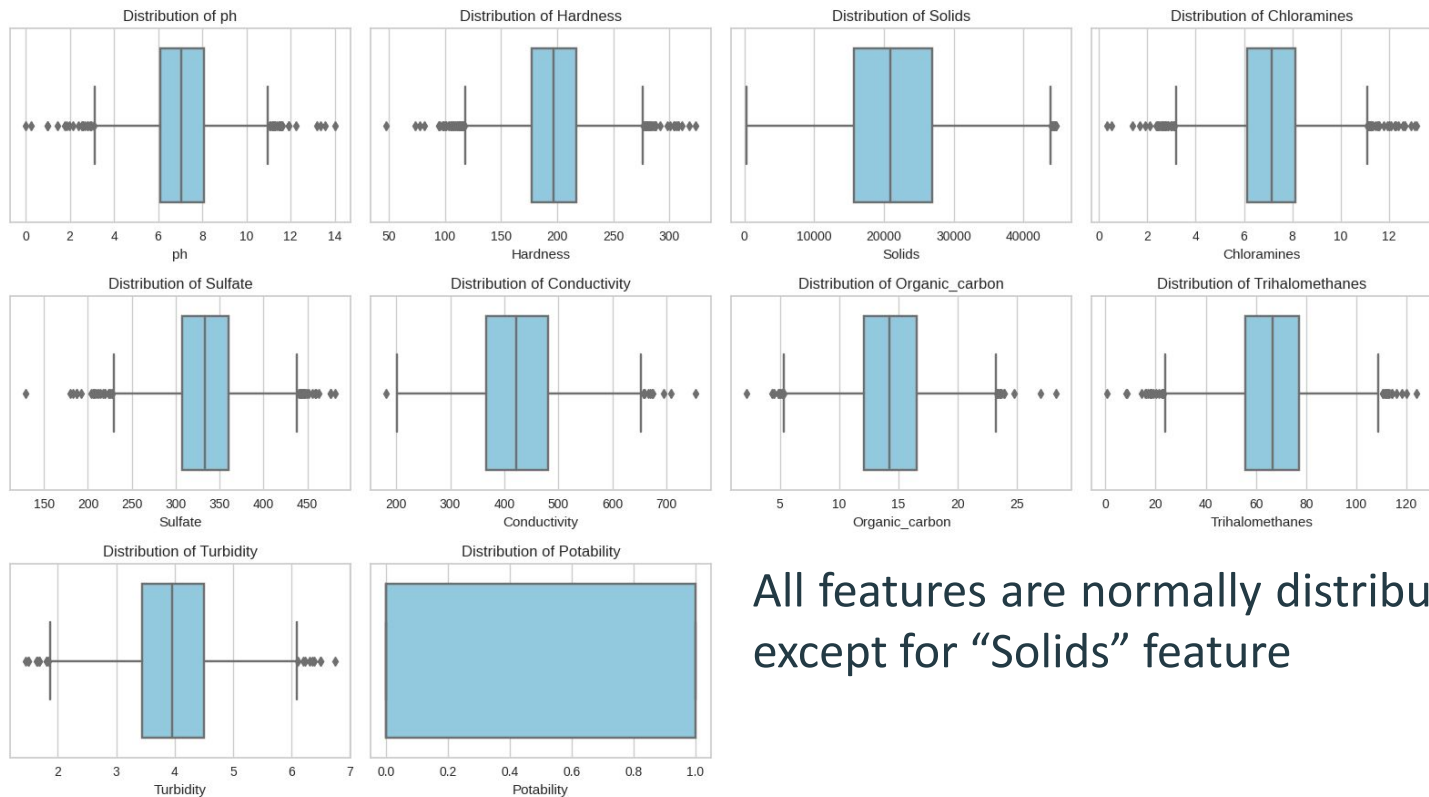
Distribution with Histplot



All features are normally distributed except for “Solids” feature



Distribution with Boxplot




All features are normally distributed except for “Solids” feature



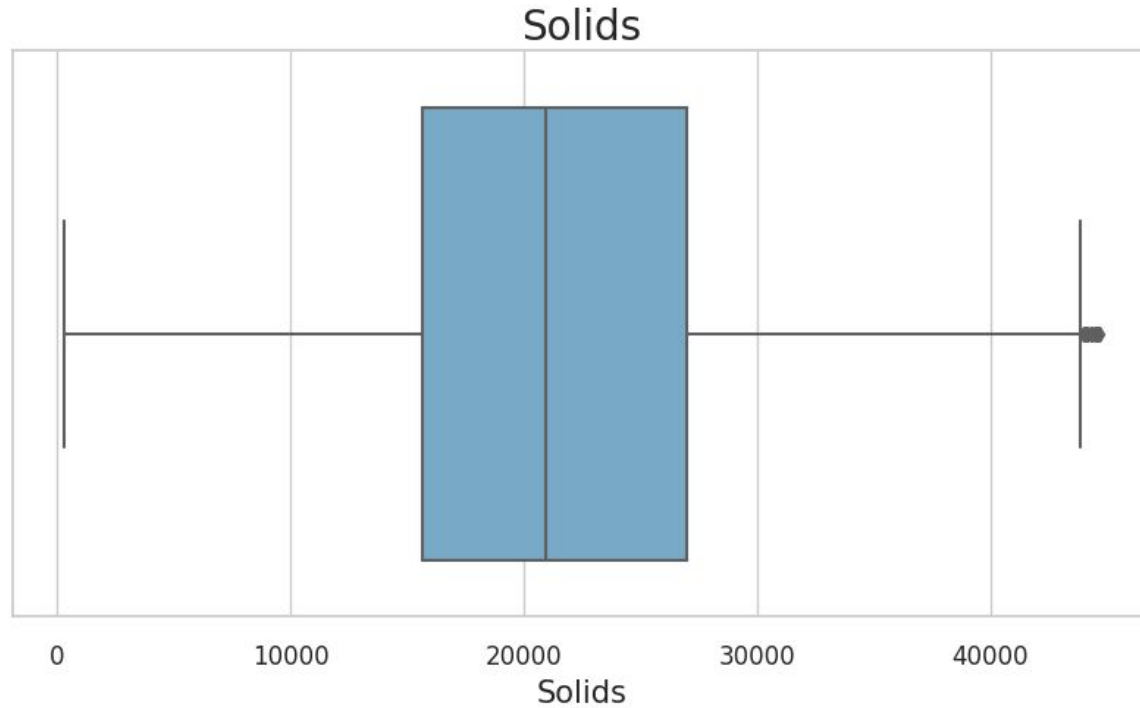
Correlation Matrix

All features seems to have low correlation with Potability.
We will explore with Non-Linear model

Correlation Matrix Heatmap

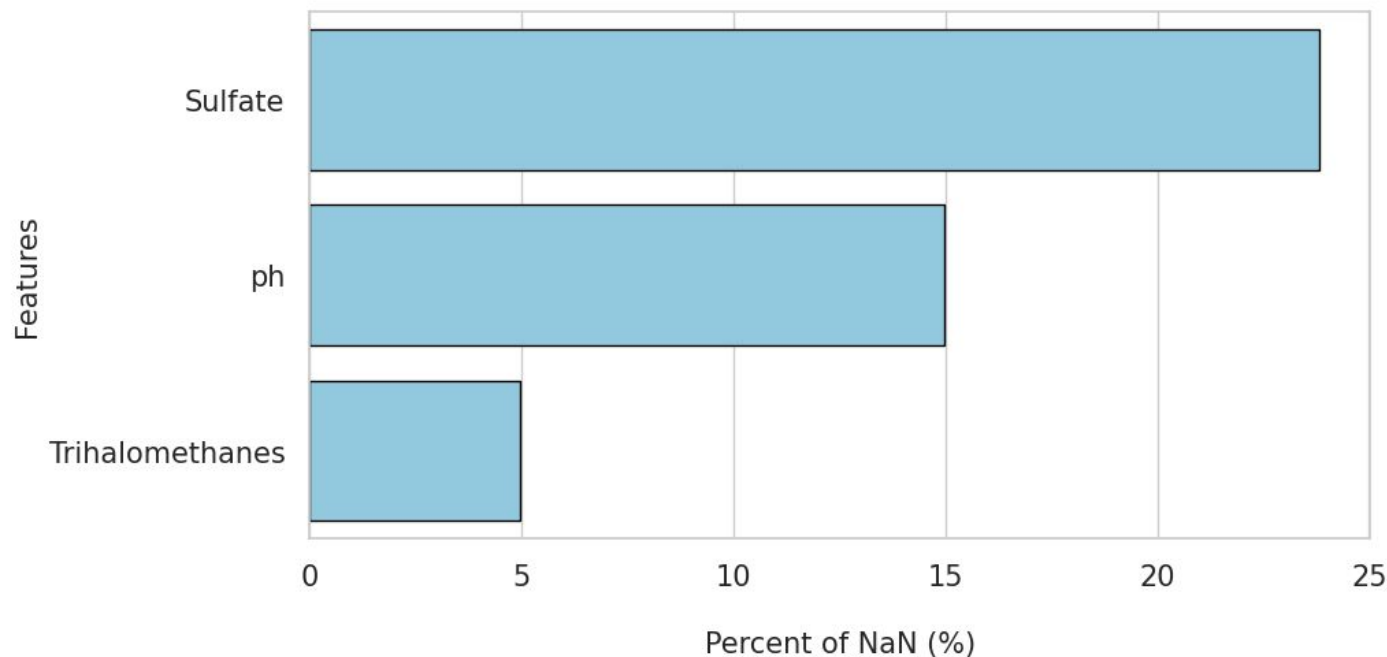


Cleaning Outliers


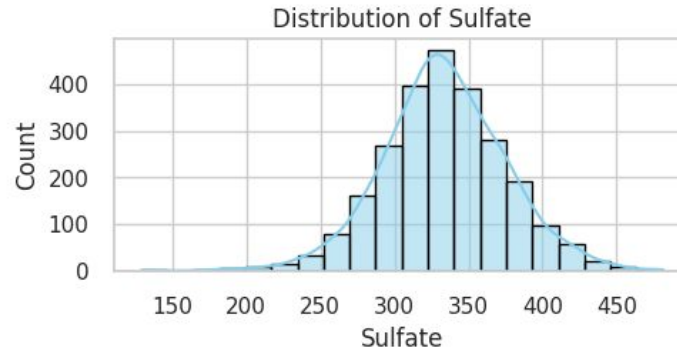



Check NaN Values

Percent of NaN per column of the dataset



Check NaN Values









ph and Trihalomethanes has normal distribution = we use mean or median
Sulfate has left skew = we use mean

To simplify, all feature will use mean
SimpleImputer



Compare Before after Imputer

Distribution is not affected after SimpleImputer



Split data set into training and testing

```
[ ] 1 # Generate training and test datasets of dependent and independent variables
    2 X_train, X_test, y_train, y_test = train_test_split(predictor, response,
    3                                                       stratify=response,
    4                                                       test_size = 0.2,
    5                                                       random_state = 0)
    6
    7 #to resolve any class imbalance - use stratify parameter
    8 print(" X_train dataset: ", X_train.shape)
    9 print(" y_train dataset: ", y_train.shape)
   10 print(" X_test dataset: ", X_test.shape)
   11 print(" y_test dataset: ", y_test.shape)
```

```
X_train dataset: (2620, 9)
y_train dataset: (2620,)
X_test dataset: (656, 9)
y_test dataset: (656,)
```

Generate training and testing datasets from **df_imputed**. First, define a predictor (X) and response (y) variable. And then split with the proportion are 80% for training and 20% for testing.



Handling Imbalanced Data with SMOTE

Before OverSampling, counts of label '1': 1022

Before OverSampling, counts of label '0': 1598

```
5 sm = SMOTE(random_state=0)
6 X_train_resampled, y_train_resampled = sm.fit_resample(X_train, y_train)
```

After OverSampling, the shape of train_X: (3196, 9)

After OverSampling, the shape of train_y: (3196,)

After OverSampling, counts of label '1' y_train_resampled: 1598

After OverSampling, counts of label '0' y_train_resampled: 1598

After OverSampling, counts of label '1' y_test: 256

After OverSampling, counts of label '0' y_test: 400



Feature Scaling

```
[ ] 1 # Feature Scaling
    2
    3 sc_X = StandardScaler()
    4 X_train2 = pd.DataFrame(sc_X.fit_transform(X_train_resampled))
    5 X_train2.columns = X_train_resampled.columns.values
    6 X_train2.index = X_train_resampled.index.values
    7 X_train_resampled = X_train2
    8
    9 X_test2 = pd.DataFrame(sc_X.transform(X_test))
   10 X_test2.columns = X_test.columns.values
   11 X_test2.index = X_test.index.values
   12 X_test = X_test2
```

Feature scaling using **StandardScaler()**, it makes mean = 0 and scales the data to unit variance.



Model

To enhance the accuracy and efficacy of water quality classification, this analysis uses the **Scikit-learn library** to implement six (6) Classifiers as below:

1. Support Vector Classifier (SVC)
2. Multi-Layer Perceptron Classifier (MPL)
3. XGBoost Classifier (XGB)
4. Decision Tree Classifier (DT)
5. Random Forest Classifier (RF)
6. K-Neighbor Classifier (KNN)

Which one is the best?



Model

Support Vector Classifier (SVC)

```
1 # Initialization MLP
2 clf_svc = SVC(probability=True, random_state=0)
3
4 # Model training
5 clf_svc.fit(X_train_resampled, y_train_resampled)
6
7 # Predicting the Test set results
8 y_pred = clf_svc.predict(X_test)
9
10 #Evaluate results
11 acc = accuracy_score(y_test, y_pred)
12 prec = precision_score(y_test, y_pred)
13 rec = recall_score(y_test, y_pred)
14 f1 = f1_score(y_test, y_pred)
15
16 model_results = pd.DataFrame([['MLP', acc, prec, rec, f1]],
17                               columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
18
19 results = model_results.sort_values(['Precision', 'Recall'], ascending = False)
20 print(results)
```

	Model	Accuracy	Precision	Recall	F1 Score
0	MLP	0.61128	0.501845	0.53125	0.516129

```
[112] 1 # Parameters to be tuned
2     param_grid = {
3         'C': [0.1, 1, 10, 100],
4         'gamma': ['scale', 'auto', 0.001, 0.01, 0.1],
5         'kernel': ['linear', 'rbf', 'poly', 'sigmoid'],
6         'degree': [2, 3, 4], # Relevant only for 'poly' kernel
7     }
8
9     svc1 = SVC()
10
11 # Create Grid Search
12 grid_search1 = GridSearchCV(svc1, param_grid, n_jobs=-1, cv=5, verbose=2)
13 grid_search1.fit(X_train_resampled, y_train_resampled)
14 print("Best parameters found: ", grid_search1.best_params_)
```

Fitting 5 folds for each of 240 candidates, totalling 1200 fits
Best parameters found: {'C': 10, 'degree': 2, 'gamma': 'auto', 'kernel': 'rbf'}

```
1 #save best parameter as best_mlp
2 best_svc = grid_search1.best_estimator_
3
4 # Predicting the test set results using best parameter
5 y_pred = best_svc.predict(X_test)
6
7 #Evaluate results
8 acc = accuracy_score(y_test, y_pred)
9 prec = precision_score(y_test, y_pred)
10 rec = recall_score(y_test, y_pred)
11 f1 = f1_score(y_test, y_pred)
12
13 #Save model performance to model_list
14 model_list.append(best_svc.__class__.__name__)
15 accuracy_list.append(acc)
16 precision_list.append(prec)
17 recall_list.append(rec)
18 f1_score_list.append(f1)
19
20 model_results = pd.DataFrame([['SVC', acc, prec, rec, f1]],
21                               columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
22
23 results_svc = model_results.append(model_results, ignore_index=True)
24 model_results
```

Model	Accuracy	Precision	Recall	F1 Score
SVC	0.625	0.517606	0.574219	0.544444



Model

Multi-Layer Perceptron Classifier (MPL)

```
1 # Initialization MLP
2 mlp = MLPClassifier(hidden_layer_sizes=(50), max_iter=1000, random_state=0)
3
4 # Model training
5 mlp.fit(X_train_resampled, y_train_resampled)
6
7 # Predicting the Test set results
8 y_pred = mlp.predict(X_test)
9
10 #Evaluate results
11 acc = accuracy_score(y_test, y_pred)
12 prec = precision_score(y_test, y_pred)
13 rec = recall_score(y_test, y_pred)
14 f1 = f1_score(y_test, y_pred)
15
16 model_results = pd.DataFrame([['MLP', acc, prec, rec, f1]],
17                               columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
18
19 results = model_results.sort_values(["Precision", "Recall"], ascending = False)
20 print(results)
```

	Model	Accuracy	Precision	Recall	F1 Score
0	MLP	0.606707	0.49635	0.53125	0.513208

```
1 # Parameters to be tuned
2 param_grid = {
3     'hidden_layer_sizes': [(50,)],
4     'activation': ['tanh', 'relu'],
5     'solver': ['sgd', 'adam'],
6     'alpha': [0.0001, 0.05],
7     'learning_rate': ['constant', 'adaptive'],
8 }
9
10 mlp2 = MLPClassifier(max_iter=1000)
11
12 # Create Grid Search
13 grid_search = GridSearchCV(mlp, param_grid, n_jobs=-1, cv=5, verbose=2)
14 grid_search.fit(X_train_resampled, y_train_resampled)
15 print("Best parameters found: ", grid_search.best_params_)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits
Best parameters found: {'activation': 'relu', 'alpha': 0.05, 'hidden_layer_sizes': (50,), 'learning_rate': 'constant', 'solver': 'adam'}

```
[110] 1 #save best parameter as best_mlp
2     best_mlp = grid_search.best_estimator_
3
4 # Predicting the test set results using best parameter
5 y_pred = best_mlp.predict(X_test)
6
7 #Evaluate results
8 acc = accuracy_score(y_test, y_pred)
9 prec = precision_score(y_test, y_pred)
10 rec = recall_score(y_test, y_pred)
11 f1 = f1_score(y_test, y_pred)
12
13 #Save model performance to model_list
14 model_list.append(best_mlp.__class__.__name__)
15 accuracy_list.append(acc)
16 precision_list.append(prec)
17 recall_list.append(rec)
18 f1_score_list.append(f1)
19
20 model_results = []
21 model_results = pd.DataFrame([['MLP', acc, prec, rec, f1]],
22                               columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
23
24 results_mlp = model_results.append(model_results, ignore_index=True)
25 model_results
```

	Model	Accuracy	Precision	Recall	F1 Score
	MLP	0.612805	0.503788	0.519531	0.511538



Model

XG Boost Classifier :

It's an extension of the traditional gradient boosting algorithm and is known for its high performance and efficiency

```
1 from xgboost import XGBClassifier
2 from sklearn.model_selection import RandomizedSearchCV
3 from sklearn.metrics import classification_report, accuracy_score
4 from sklearn.model_selection import RepeatedStratifiedKFold
5
6 # Create an XGBoost model
7 XGB_model = XGBClassifier()
8
9 # Define hyperparameter search space
10 param_dist = {
11     "learning_rate": [0.05, 0.10, 0.15, 0.20, 0.25, 0.30],
12     "max_depth": [3, 4, 5, 6, 8, 10, 12, 15],
13     "min_child_weight": [1, 3, 5, 7],
14     "gamma": [0.0, 0.1, 0.2, 0.3, 0.4],
15     "colsample_bytree": [0.3, 0.4, 0.5, 0.7]
16 }
17
18 # Create cross-validation strategy
19 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
20
21 # Perform hyperparameter tuning with RandomizedSearchCV
22 XGB = RandomizedSearchCV(XGB_model, param_dist, n_iter=50, scoring='roc_auc', n_jobs=-1, cv=cv, random_state=1)
23 XGB.fit(X_train_resampled, y_train_resampled)
24
25 # Predict on the test data
26 y_pred = XGB.predict(X_test)
27
28 # Evaluate results
29 acc = accuracy_score(y_test, y_pred)
30 prec = precision_score(y_test, y_pred)
31 rec = recall_score(y_test, y_pred)
32 f1 = f1_score(y_test, y_pred)
```

```
34 # Display classification report and accuracy
35 print(classification_report(y_test, y_pred))
36 accuracy = accuracy_score(y_test, y_pred)
37 print(f'Accuracy: {accuracy}')
38
39 # Get the best estimator from RandomizedSearchCV
40 best_XGB = XGB.best_estimator_
41
42 # Save model performance to model_list
43 model_list.append([best_XGB.__class__.__name__,
44 accuracy_list.append(acc)
45 precision_list.append(prec)
46 recall_list.append(rec)
47 f1_score_list.append(f1)
```

	precision	recall	f1-score	support
0.0	0.69	0.67	0.68	400
1.0	0.51	0.53	0.52	256
accuracy			0.62	656
macro avg	0.60	0.60	0.60	656
weighted avg	0.62	0.62	0.62	656



Model

Decision Tree

A decision tree is a popular supervised machine learning algorithm used for both classification and regression tasks..

```
1 #from sklearn.tree import DecisionTreeClassifier
2 from sklearn.model_selection import RandomizedSearchCV
3 from scipy.stats import randint
4 from sklearn.metrics import classification_report
5 from sklearn.model_selection import RepeatedStratifiedKFold
6
7 # Kemudian, Anda dapat membuat objek RepeatedStratifiedKFold seperti ini:
8 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
9
10
11
12 # Create a Decision Tree model
13 DT_model = DecisionTreeClassifier()
14
15 # Define hyperparameter search space
16 param_dist = {
17     "max_depth": [3, None],
18     "max_features": randint(1, X_train_resampled.shape[1]),
19     "min_samples_leaf": randint(1, X_train_resampled.shape[0]),
20     "criterion": ["gini", "entropy"]
21 }
22
23 # Create cross-validation strategy
24 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
25
26 # Perform hyperparameter tuning with RandomizedSearchCV
27 DT = RandomizedSearchCV(DT_model, param_dist, n_iter=50, scoring='roc_auc', n_jobs=-1, cv=cv, random_state=1)
28 DT.fit(X_train_resampled, y_train_resampled)
```

```
30 # Get the best estimator from RandomizedSearchCV
31 best_DT = DT.best_estimator_
32
33 # Train the best Decision Tree model on the training data
34 best_DT.fit(X_train_resampled, y_train_resampled)
35
36 # Predict on the test data
37 pred = best_DT.predict(X_test)
38 pred_prob = best_DT.predict_proba(X_test)
39
40 # Evaluate results
41 acc = accuracy_score(y_test, pred)
42 prec = precision_score(y_test, pred)
43 rec = recall_score(y_test, pred)
44 f1 = f1_score(y_test, pred)
45
46 # Save model performance to model_list
47 model_list.append(best_DT.__class__.__name__)
48 accuracy_list.append(acc)
49 precision_list.append(prec)
50 recall_list.append(rec)
51 f1_score_list.append(f1)
52
53 # Display classification report
54 print(classification_report(y_test, pred))
55
```

	precision	recall	f1-score	support
0.0	0.63	0.53	0.58	400
1.0	0.41	0.52	0.46	256
accuracy			0.52	656
macro avg	0.52	0.52	0.52	656
weighted avg	0.55	0.52	0.53	656

Model

Random Forest :

Random Forests are made out of decision trees composed of different bootstrapped data.

```
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.model_selection import RandomizedSearchCV, RepeatedStratifiedKFold
5 from scipy.stats import randint
6 from sklearn.metrics import classification_report
7
8 # Create a Random Forest model
9 RF_model = RandomForestClassifier()
10
11 # Define hyperparameter search space for Random Forest
12 RF_param_dist = {
13     "n_estimators": [100, 200, 300],
14     "max_depth": [None, 10, 20, 30],
15     "min_samples_split": [2, 5, 10],
16     "min_samples_leaf": [1, 2, 4],
17     "criterion": ["gini", "entropy"]
18 }
19
20 # Create cross-validation strategy
21 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
22
23 # Perform hyperparameter tuning with RandomizedSearchCV for Random Forest
24 RF = RandomizedSearchCV(RF_model, RF_param_dist, n_iter=50, scoring='roc_auc', n_jobs=-1, cv=cv, random_state=1)
25 RF.fit(X_train_resampled, y_train_resampled)
26
27 # Get the best estimator from RandomizedSearchCV for Random Forest
28 best_RF = RF.best_estimator_
29
```

```
1 # Train the best Random Forest model on the training data
2 best_RF.fit(X_train_resampled, y_train_resampled)
3
4 # Predict on the test data for Random Forest
5 pred_RF = best_RF.predict(X_test)
6 pred_prob_RF = best_RF.predict_proba(X_test)
7
8 # Evaluate results
9 acc = accuracy_score(y_test, y_pred)
10 prec = precision_score(y_test, y_pred)
11 rec = recall_score(y_test, y_pred)
12 f1 = f1_score(y_test, y_pred)
13
14 # Save model performance to model_list
15 model_list.append(best_RF.__class__.__name__)
16 accuracy_list.append(acc)
17 precision_list.append(prec)
18 recall_list.append(rec)
19 f1_score_list.append(f1)
20
21 # Display classification report for Random Forest
22 print("Classification Report for Random Forest:")
23 print(classification_report(y_test, pred_RF))

```

Classification Report for Random Forest:					
	precision	recall	f1-score	support	
	0.0	0.71	0.72	0.71	400
	1.0	0.55	0.53	0.54	256
accuracy				0.64	656
macro avg	0.63	0.62	0.62		656
weighted avg	0.64	0.64	0.64		656



Model

K-Nearest Neighbour :

K-NN makes predictions by finding the **K-nearest data** points in the training dataset to a given data point and then classifying or regressing based on the majority or average of the neighbors' labels or values.

```
1 # K-Nearest Neighbors (KNN) model
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.model_selection import RandomizedSearchCV, RepeatedStratifiedKFold
4 from scipy.stats import randint
5 from sklearn.metrics import classification_report
6
7 # Create a K-Nearest Neighbors (KNN) model
8 KNN_model = KNeighborsClassifier()
9
10 # Define hyperparameter search space for KNN
11 KNN_param_dist = {
12     "n_neighbors": randint(1, 10),
13     "weights": ['uniform', 'distance'],
14     "algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute']
15 }
16
17 # Create cross-validation strategy
18 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
19
20 # Perform hyperparameter tuning with RandomizedSearchCV for KNN
21 KNN = RandomizedSearchCV(KNN_model, KNN_param_dist, n_iter=50, scoring='roc_auc', n_jobs=-1, cv=cv, random_state=1)
22 KNN.fit(X_train_resampled, y_train_resampled)
23
24 # Get the best estimator from RandomizedSearchCV for KNN
25 best_KNN = KNN.best_estimator_
26
```

```
1 # Train the best KNN model on the training data
2 best_KNN.fit(X_train_resampled, y_train_resampled)
3
4 # Predict on the test data for KNN
5 pred_KNN = best_KNN.predict(X_test)
6 pred_prob_KNN = best_KNN.predict_proba(X_test)
7
8 # Evaluate results
9 acc = accuracy_score(y_test, pred_KNN)
10 prec = precision_score(y_test, pred_KNN)
11 rec = recall_score(y_test, pred_KNN)
12 f1 = f1_score(y_test, pred_KNN)
13
14 # Save model performance to model_list
15 model_list.append(best_KNN.__class__.__name__)
16 accuracy_list.append(acc)
17 precision_list.append(prec)
18 recall_list.append(rec)
19 f1_Score_list.append(f1)
20
21 # Display classification report for KNN
22 print("Classification Report for K-Nearest Neighbors:")
23 print(classification_report(y_test, pred_KNN))
24
```

Classification Report for K-Nearest Neighbors:					
	precision	recall	f1-score	support	
	0.0	0.69	0.59	0.63	400
	1.0	0.47	0.58	0.52	256
accuracy				0.59	656
macro avg	0.58	0.58	0.58		656
weighted avg	0.60	0.59	0.59		656



Model Evaluation

Accuracy Model

	Model	Accuracy	Precision	Recall	F1 Score
0	MLPClassifier	0.612805	0.503788	0.519531	0.511538
1	SVC	0.625000	0.517606	0.574219	0.544444
2	RandomForestClassifier	0.641768	0.543210	0.515625	0.529058
3	KNeighborsClassifier	0.585366	0.474522	0.582031	0.522807
4	DecisionTreeClassifier	0.591463	0.476923	0.484375	0.480620
5	XGBClassifier	0.615854	0.507519	0.527344	0.517241

"Random Forest Classifier (RF) with accuracy of 0.64 is the best one"




Model Evaluation

Confusion-Matrix

		Actual Values	
		1 (Positive)	0 (Negative)
Predicted Values	1 (Positive)	TP (True Positive)	FP (False Positive) <small>Type I Error</small>
	0 (Negative)	FN (False Negative) <small>Type II Error</small>	TN (True Negative)

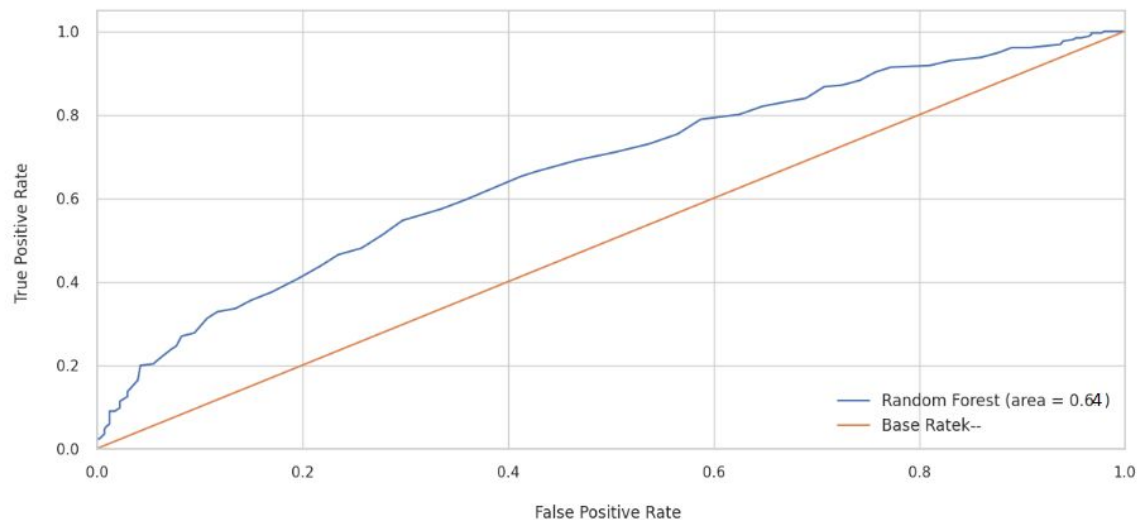
Confusion Matrix



Model Evaluation

Evaluate the model using ROC Graph

ROC Graph



Model Deployment

Streamlit

110

Water Potability Prediction App

Battleship Team

Home

Water Potability Prediction App

This app will be used to predict water potability

Data Source

- https://raw.githubusercontent.com/rene-gith/water-potability/main/water_potability.csv

App Content

- Exploratory Data Analysis
- Machine Learning Section

Import Data

Chloramines concentration

0 - +

Sulfate concentration

129 - +

Electrical conductivity

181 - +

Organic carbon content

2 - +

Trihalomethanes concentration

0 - +

Turbidity level

1 - +

Prediction Result

Open in Browser



C. Conclusion

"Random Forest Classifier (RF) with accuracy of 0.64 is the best one to other models in accuracy and robustness, making it the optimal choice for classification."

